


OBD SuperStar Agent

Architecture Overview

A Next.js (React) frontend with a Python FastAPI backend powering a multi-agent pipeline. The system takes three inputs -- **product documentation**, **country/telco**, and **target audience** -- and outputs fully produced OBD audio files ready for deployment.



Agent Pipeline (6 Agents)

Agent 1: Product Analyzer

- **Input:** Product documentation (PDF/PPT/text), pricing info
- **Role:** Reads and fully understands the product -- features, pricing, USPs, target use cases
- **Output:** Structured product brief (JSON) with key selling points, pricing tiers, and value propositions

Agent 2: Market Researcher

- **Input:** Country name, Telco name, product brief from Agent 1

- **Role:** Acts as a local market analyst. Researches the country's demographics, mobile usage patterns, cultural nuances, current affairs, local trends, feature phone vs smartphone split, language preferences
- **Output:** Market analysis report with target audience psyche, cultural hooks, current affairs angles, and language recommendations

Agent 3: Script Writer (Hook + Body + CTA)

- **Input:** Product brief + Market analysis
- **Role:** Acts as the best copywriter in that country who understands local culture, nuances, and psychology. Creates scripts under 30 seconds with:
 - **Hook** (first 5 seconds): Must prevent call disconnect. Culturally relevant, attention-grabbing
 - **Body:** Product pitch delivered compellingly so recipients want to subscribe
 - **CTA:** DTMF-based call to action
 - **Fallback 1:** If no DTMF pressed -- urgency-based follow-up CTA
 - **Fallback 2:** Third attempt using psychological techniques (scarcity, loss aversion, social proof)
 - **Polite closure:** If still no response
- **Output:** 5 complete script sets, each with main flow + 2 fallback variants. Scripts include ElevenLabs V3 audio tags ([excited], [whispers], [pause], etc.) baked in
- **Key constraint:** Each script must be under 30 seconds when spoken

Agent 4: Evaluation Panel

- **Input:** All 5 script sets from Agent 3
- **Role:** A panel of 10 AI evaluators, each with a distinct persona:
 1. Award-winning radio jingle writer
 2. Consumer psychologist
 3. Behavioral economist
 4. Top copywriter #1 (direct response specialist)
 5. Top copywriter #2 (brand storytelling)
 6. Top copywriter #3 (local market expert)
 7. Dark psychology / persuasion expert
 8. Telco marketing director
 9. UX/voice interface designer
 10. Local cultural consultant
- **Process:** Each evaluator scores and critiques. Consensus feedback is sent back to Agent 3 for revision. One iteration loop.
- **Output:** Final approved script sets with scores and rationale

Agent 5: Voice Selector

- **Input:** Final scripts + country/language info
- **Role:** Queries ElevenLabs API (GET /v2/voices) to find voices matching the target country/language. Evaluates voice characteristics (gender, age, tone, accent) against the

- script requirements. Recommends optimal voice with API parameters (stability, similarity_boost, style, speed).
- **Output:** Voice ID, voice settings JSON, and rationale for selection

Agent 6: Audio Producer

- **Input:** Final scripts (with V3 audio tags) + Voice selection
- **Role:** Calls ElevenLabs TTS API (`POST /v1/text-to-speech/:voice_id`) using the `eleven_v3` model. Generates audio for all script variants. Uses V3 audio tags for expressiveness: `[excited]`, `[whispers]`, `[pause]`, `[laughs]`, `[sigh]`, `[curious]`, capitalization for emphasis, ellipses for pacing.
- **Output:** MP3 audio files for each script variant, downloadable from the UI

Tech Stack

Frontend: Next.js + Tailwind CSS + shadcn/ui

- `frontend/` directory
- Modern, clean UI with a step-by-step wizard
- Real-time progress tracking as agents execute
- Audio player for previewing generated recordings
- Download functionality for final audio files

Backend: Python FastAPI

- `backend/` directory
- Agent orchestration engine
- LLM integration supporting both **Claude** (Anthropic) and **OpenAI GPT-4** (configurable)
- ElevenLabs API integration
- WebSocket for real-time progress updates to frontend

Key Dependencies

- **Python:** fastapi, uvicorn, anthropic, openai, elevenlabs, python-multipart, websockets, pydantic
- **Node:** next, react, tailwindcss, shadcn/ui, lucide-react

Project Structure

```

OBD_SuperStarAgent/
  backend/
    main.py          # FastAPI app, routes, WebSocket
    agents/
      __init__.py
      base.py        # Base agent class with LLM abstraction
  
```

```

product_analyzer.py      # Agent 1
market_researcher.py     # Agent 2
script_writer.py         # Agent 3
eval_panel.py            # Agent 4
voice_selector.py        # Agent 5
audio_producer.py       # Agent 6
orchestrator.py          # Pipeline orchestrator
config.py                # API keys, model config
requirements.txt

frontend/
  app/
    page.tsx              # Main wizard UI
    layout.tsx
    globals.css
  components/
    ProductUpload.tsx      # Step 1: Upload product docs
    CountryTelcoSelect.tsx # Step 2: Select country/telco
    PipelineProgress.tsx   # Step 3: Real-time agent progress
    ScriptReview.tsx        # Step 4: Review generated scripts
    AudioPlayer.tsx         # Step 5: Listen and download audio
  package.json
  tailwind.config.ts
.env                      # API keys (ANTHROPIC_API_KEY, OPENAI_API_KEY,
ELEVENLABS_API_KEY)
README.md

```

Key Design Decisions

- **LLM Abstraction:** Base agent class supports both Claude and OpenAI. Each agent call specifies which provider to use via config. Default: Claude for creative tasks (script writing, evaluation), OpenAI as fallback.
- **Feedback Loop:** Agent 4 (Eval Panel) sends structured feedback to Agent 3 (Script Writer) for exactly one revision cycle to balance quality vs. speed.
- **V3 Audio Tags:** Agent 3 outputs scripts pre-formatted with ElevenLabs V3 audio tags ([excited], [whispers], [pause], etc.) so Agent 6 can pass them directly to the API.
- **30-Second Constraint:** Agent 3 estimates word count (~150 words = 60 seconds, ~75 words = 30 seconds) and enforces the limit. Agent 6 validates actual audio duration.
- **Streaming Progress:** WebSocket connection from frontend to backend shows real-time agent status (which agent is running, intermediate outputs).